

Feed_Forward_Network_using_Pytorch_with_MNIST

April 18, 2022

1 Implementing DNN and CNN on MNIST dataset

Implemented 2 and 3 hidden layer DNN and a single convolutional layer CNN for classification of digits on MNIST dataset

Libraries used:

- 1) PyTorch - for implementing DNN and CNN
- 2) Matplotlib - to plot the graph

```
[1]: import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
```

```
[2]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

Setting the hyper-parameters

```
[3]: input_size = 784 # 28x28
hidden_size = 512
num_classes = 10
num_epochs = 5
batch_size = 100
learning_rate = 0.001
```

Importing the MNIST dataset

```
[4]: train_dataset = torchvision.datasets.MNIST(root='./data',
                                                train=True,
                                                transform=transforms.ToTensor(),
                                                download=True)

test_dataset = torchvision.datasets.MNIST(root='./data',
                                           train=False,
                                           transform=transforms.ToTensor())
```

```
[5]: print(train_dataset)
print(test_dataset)
```

```
Dataset MNIST
  Number of datapoints: 60000
  Root location: ./data
  Split: Train
  StandardTransform
Transform: ToTensor()
Dataset MNIST
  Number of datapoints: 10000
  Root location: ./data
  Split: Test
  StandardTransform
Transform: ToTensor()
```

Loading the datasets into mini-batches of size 100 samples

```
[6]: train_loader = torch.utils.data.DataLoader(dataset=train_dataset,
                                                batch_size=batch_size,
                                                shuffle=True)

test_loader = torch.utils.data.DataLoader(dataset=test_dataset,
                                           batch_size=batch_size,
                                           shuffle=False)
```

1.0.1 Question 2) a) Class DeepNeuralNet

Implements the DNN with 2 hidden layers and
ReLU hidden activation on both the hidden layers
with hidden neuron size = 512

```
[7]: class DeepNeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size, num_classes):
        super(DeepNeuralNet, self).__init__()
        self.input_size = input_size
        self.l1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.l2 = nn.Linear(hidden_size, hidden_size)
        self.l3 = nn.Linear(hidden_size, num_classes)

    def forward(self, x):
        out = self.l1(x)
        out = self.relu(out)
        out = self.l2(out)
        out = self.relu(out)
        out = self.l3(out)
        return out

model = DeepNeuralNet(input_size, hidden_size, num_classes).to(device)
```

Using Cross-Entropy Loss and Adam Optimizer

```
[8]: criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
```

1.0.2 Function train(model, optimizer, type)

Trains the given model using the MNIST dataset in mini-batches and
Plots the Loss curve for each step of SGD.

```
[9]: def train(model, optimizer, type = "DNN"):
    n_total_steps = len(train_loader)

    loss_plot = []
    for epoch in range(num_epochs):
        for i, (images, labels) in enumerate(train_loader):
            # origin shape: [100, 1, 28, 28]
            # resized: [100, 784]

            if type == "DNN":
                images = images.reshape(-1, 28*28).to(device)
            else:
                images = images.to(device)
            labels = labels.to(device)

            # Forward pass
            outputs = model(images)
            loss = criterion(outputs, labels)

            # Backward and optimize
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            if (i+1) % 100 == 0:
                print (f'Epoch [{epoch+1}/{num_epochs}], Step [{i+1}/
↪{n_total_steps}], Loss: {loss.item():.4f}')
                loss_plot.append(loss.item())

    plt.plot([i for i in range(num_epochs * 6)], loss_plot, marker="o",
↪color='b')
    plt.xlabel('Epoch - Mini Batch')
    plt.ylabel('Loss')
    plt.title('Loss curve')
    plt.show()
```

1.0.3 Function predict(model, type)

Predicts the output of the data point using the trained model and calculates the test accuracy of the model.

```
[10]: def predict(model, type = "DNN"):
    with torch.no_grad():
        n_correct = 0
        n_samples = 0
        for images, labels in test_loader:
            if type == "DNN":
                images = images.reshape(-1, 28*28).to(device)
            else:
                images = images.to(device)

            labels = labels.to(device)

            outputs = model(images)

            # max returns (value ,index)
            _, predicted = torch.max(outputs.data, 1)
            n_samples += labels.size(0)
            n_correct += (predicted == labels).sum().item()

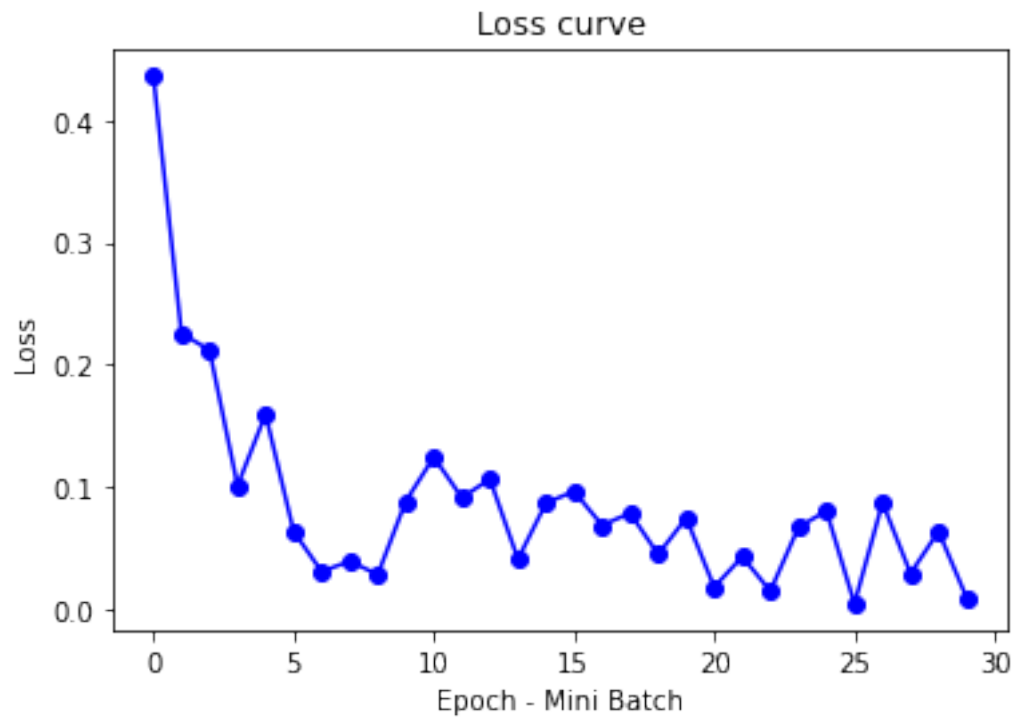
        acc = 100.0 * n_correct / n_samples
        print(f'#### Test Accuracy of the network on 10000 test images: {acc} %\n
        ↪####')
```

Training the DNN with 2 hidden layer

```
[11]: train(model, optimizer)
```

```
Epoch [1/5], Step [100/600], Loss: 0.4363
Epoch [1/5], Step [200/600], Loss: 0.2255
Epoch [1/5], Step [300/600], Loss: 0.2119
Epoch [1/5], Step [400/600], Loss: 0.1008
Epoch [1/5], Step [500/600], Loss: 0.1594
Epoch [1/5], Step [600/600], Loss: 0.0634
Epoch [2/5], Step [100/600], Loss: 0.0304
Epoch [2/5], Step [200/600], Loss: 0.0390
Epoch [2/5], Step [300/600], Loss: 0.0280
Epoch [2/5], Step [400/600], Loss: 0.0881
Epoch [2/5], Step [500/600], Loss: 0.1242
Epoch [2/5], Step [600/600], Loss: 0.0910
Epoch [3/5], Step [100/600], Loss: 0.1062
Epoch [3/5], Step [200/600], Loss: 0.0406
Epoch [3/5], Step [300/600], Loss: 0.0866
Epoch [3/5], Step [400/600], Loss: 0.0962
Epoch [3/5], Step [500/600], Loss: 0.0684
```

Epoch [3/5], Step [600/600], Loss: 0.0783
 Epoch [4/5], Step [100/600], Loss: 0.0445
 Epoch [4/5], Step [200/600], Loss: 0.0743
 Epoch [4/5], Step [300/600], Loss: 0.0177
 Epoch [4/5], Step [400/600], Loss: 0.0426
 Epoch [4/5], Step [500/600], Loss: 0.0144
 Epoch [4/5], Step [600/600], Loss: 0.0666
 Epoch [5/5], Step [100/600], Loss: 0.0802
 Epoch [5/5], Step [200/600], Loss: 0.0042
 Epoch [5/5], Step [300/600], Loss: 0.0869
 Epoch [5/5], Step [400/600], Loss: 0.0289
 Epoch [5/5], Step [500/600], Loss: 0.0634
 Epoch [5/5], Step [600/600], Loss: 0.0090



Calculating the Test Accuracy of the trained DNN with 2 hidden layer Test Accuracy
 = 97.89%

[12]: `predict(model)`

Test Accuracy of the network on 10000 test images: 97.89 %

1.0.4 Class DeepNeuralNet_Three

Implements the DNN with 3 hidden layers and
ReLU hidden activation on all the hidden layers
with hidden neuron size = 512

```
[13]: class DeepNeuralNet_Three(nn.Module):
    def __init__(self, input_size, hidden_size, num_classes):
        super(DeepNeuralNet_Three, self).__init__()
        self.input_size = input_size
        self.l1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.l2 = nn.Linear(hidden_size, hidden_size)
        self.l3 = nn.Linear(hidden_size, hidden_size)
        self.l4 = nn.Linear(hidden_size, num_classes)

    def forward(self, x):
        out = self.l1(x)
        out = self.relu(out)
        out = self.l2(out)
        out = self.relu(out)
        out = self.l3(out)
        out = self.relu(out)
        out = self.l4(out)
        return out

model = DeepNeuralNet_Three(input_size, hidden_size, num_classes).to(device)
```


```
[14]: optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
```

Training the DNN with 3 hidden layer

```
[15]: train(model, optimizer)
```

```
Epoch [1/5], Step [100/600], Loss: 0.4178
Epoch [1/5], Step [200/600], Loss: 0.1536
Epoch [1/5], Step [300/600], Loss: 0.0940
Epoch [1/5], Step [400/600], Loss: 0.1802
Epoch [1/5], Step [500/600], Loss: 0.1378
Epoch [1/5], Step [600/600], Loss: 0.1143
Epoch [2/5], Step [100/600], Loss: 0.0962
Epoch [2/5], Step [200/600], Loss: 0.0283
Epoch [2/5], Step [300/600], Loss: 0.0458
Epoch [2/5], Step [400/600], Loss: 0.1278
Epoch [2/5], Step [500/600], Loss: 0.2246
Epoch [2/5], Step [600/600], Loss: 0.0209
Epoch [3/5], Step [100/600], Loss: 0.1053
Epoch [3/5], Step [200/600], Loss: 0.0206
Epoch [3/5], Step [300/600], Loss: 0.1545
```

Epoch [3/5], Step [400/600], Loss: 0.0325
 Epoch [3/5], Step [500/600], Loss: 0.0418
 Epoch [3/5], Step [600/600], Loss: 0.0721
 Epoch [4/5], Step [100/600], Loss: 0.0014
 Epoch [4/5], Step [200/600], Loss: 0.0858
 Epoch [4/5], Step [300/600], Loss: 0.0720
 Epoch [4/5], Step [400/600], Loss: 0.0277
 Epoch [4/5], Step [500/600], Loss: 0.1126
 Epoch [4/5], Step [600/600], Loss: 0.1445
 Epoch [5/5], Step [100/600], Loss: 0.0535
 Epoch [5/5], Step [200/600], Loss: 0.1593
 Epoch [5/5], Step [300/600], Loss: 0.0129
 Epoch [5/5], Step [400/600], Loss: 0.0810
 Epoch [5/5], Step [500/600], Loss: 0.0241
 Epoch [5/5], Step [600/600], Loss: 0.0071



Calculating the Test Accuracy of the trained DNN with 3 hidden layer Test Accuracy
 = 97.84%

```
[16]: predict(model)
```

Test Accuracy of the network on 10000 test images: 97.84 %

1.0.5 Question 2) b) Class CNN

Implements CNN with single 2D-convolutional layer
with kernel size = 3 and 128 filters followed by
two Dense layers of 256 neurons

```
[17]: class CNN(nn.Module):
    def __init__(self, filter, kernel_size, hidden_size):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(
            nn.Conv2d(
                in_channels = 1,
                out_channels = filter,
                kernel_size = kernel_size,
                stride=1,
            ),
            nn.ReLU(),
        )
        self.l2 = nn.Linear(filter * 26 * 26, hidden_size)
        self.relu = nn.ReLU()
        self.l3 = nn.Linear(hidden_size, 10)

    def forward(self, x):
        x = self.conv1(x)
        # flatten the output of conv1 to (batch_size, filter * 26 * 26)
        x = x.view(x.size(0), -1)
        x = self.l2(x)
        x = self.relu(x)
        out = self.l3(x)
        return out

model = CNN(128, 3, 256).to(device)
```

Using Cross-Entropy Loss and Adam Optimizer


```
[18]: criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
```

Training the CNN

```
[19]: train(model, optimizer, "CNN")
```

```
Epoch [1/5], Step [100/600], Loss: 0.2368
Epoch [1/5], Step [200/600], Loss: 0.2767
Epoch [1/5], Step [300/600], Loss: 0.0940
Epoch [1/5], Step [400/600], Loss: 0.1550
Epoch [1/5], Step [500/600], Loss: 0.0822
Epoch [1/5], Step [600/600], Loss: 0.0529
```


Epoch [2/5], Step [100/600], Loss: 0.0424
Epoch [2/5], Step [200/600], Loss: 0.0466
Epoch [2/5], Step [300/600], Loss: 0.1000
Epoch [2/5], Step [400/600], Loss: 0.0680
Epoch [2/5], Step [500/600], Loss: 0.1171
Epoch [2/5], Step [600/600], Loss: 0.0427
Epoch [3/5], Step [100/600], Loss: 0.0361
Epoch [3/5], Step [200/600], Loss: 0.0598
Epoch [3/5], Step [300/600], Loss: 0.0245
Epoch [3/5], Step [400/600], Loss: 0.0079
Epoch [3/5], Step [500/600], Loss: 0.0476
Epoch [3/5], Step [600/600], Loss: 0.0236
Epoch [4/5], Step [100/600], Loss: 0.0181
Epoch [4/5], Step [200/600], Loss: 0.0035
Epoch [4/5], Step [300/600], Loss: 0.0108
Epoch [4/5], Step [400/600], Loss: 0.0155
Epoch [4/5], Step [500/600], Loss: 0.0344
Epoch [4/5], Step [600/600], Loss: 0.0590
Epoch [5/5], Step [100/600], Loss: 0.0179
Epoch [5/5], Step [200/600], Loss: 0.0032
Epoch [5/5], Step [300/600], Loss: 0.0062
Epoch [5/5], Step [400/600], Loss: 0.0055
Epoch [5/5], Step [500/600], Loss: 0.0041
Epoch [5/5], Step [600/600], Loss: 0.0096



Calculating the test accuracy of CNN Test accuracy = 98.25%

```
[20]: predict(model, type = "CNN")
```

```
#### Test Accuracy of the network on 10000 test images: 98.25 % ####
```

1.1 Comparing DNN and CNN

CNN got a better test accuracy of 98.25% than both the DNN with 2 and 3 hidden layers with test accuracy of 97.89% and 97.84% respectively.

1.1.1 Question 2) c) Analysis of CNN on Different filter size and Hidden layer dimensions

For Filter sizes = [32, 64, 128, 256] and


For Hidden Layer dimension = [128, 256, 512]

```
[21]: filter = [32, 64, 128, 256]
```

```
for j in filter:
    print("\n\n##### CNN with Filters = ", j, " and Hidden layer_
↪dimension = 128 #####")
    model = CNN(j, 3, 128).to(device)
    optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
    train(model, optimizer, "CNN")
    predict(model, type = "CNN")
```

```
##### CNN with Filters = 32 and Hidden layer dimension = 128 #####
Epoch [1/5], Step [100/600], Loss: 0.4227
Epoch [1/5], Step [200/600], Loss: 0.1033
Epoch [1/5], Step [300/600], Loss: 0.0883
Epoch [1/5], Step [400/600], Loss: 0.3845
Epoch [1/5], Step [500/600], Loss: 0.1905
Epoch [1/5], Step [600/600], Loss: 0.1480
Epoch [2/5], Step [100/600], Loss: 0.0507
Epoch [2/5], Step [200/600], Loss: 0.1223
Epoch [2/5], Step [300/600], Loss: 0.0844
Epoch [2/5], Step [400/600], Loss: 0.0430
Epoch [2/5], Step [500/600], Loss: 0.0376
Epoch [2/5], Step [600/600], Loss: 0.0612
Epoch [3/5], Step [100/600], Loss: 0.0078
Epoch [3/5], Step [200/600], Loss: 0.0115
Epoch [3/5], Step [300/600], Loss: 0.0114
Epoch [3/5], Step [400/600], Loss: 0.0699
Epoch [3/5], Step [500/600], Loss: 0.0305
Epoch [3/5], Step [600/600], Loss: 0.0193
Epoch [4/5], Step [100/600], Loss: 0.0111
```

Epoch [4/5], Step [200/600], Loss: 0.0339
 Epoch [4/5], Step [300/600], Loss: 0.0154
 Epoch [4/5], Step [400/600], Loss: 0.0206
 Epoch [4/5], Step [500/600], Loss: 0.0188
 Epoch [4/5], Step [600/600], Loss: 0.0096
 Epoch [5/5], Step [100/600], Loss: 0.0089
 Epoch [5/5], Step [200/600], Loss: 0.0180
 Epoch [5/5], Step [300/600], Loss: 0.0181
 Epoch [5/5], Step [400/600], Loss: 0.0218
 Epoch [5/5], Step [500/600], Loss: 0.0057
 Epoch [5/5], Step [600/600], Loss: 0.0186




Test Accuracy of the network on 10000 test images: 98.29 %

CNN with Filters = 64 and Hidden layer dimension = 128

Epoch [1/5], Step [100/600], Loss: 0.1981
 Epoch [1/5], Step [200/600], Loss: 0.1792
 Epoch [1/5], Step [300/600], Loss: 0.2707
 Epoch [1/5], Step [400/600], Loss: 0.1944
 Epoch [1/5], Step [500/600], Loss: 0.1051
 Epoch [1/5], Step [600/600], Loss: 0.1675
 Epoch [2/5], Step [100/600], Loss: 0.1385
 Epoch [2/5], Step [200/600], Loss: 0.0372

Epoch [2/5], Step [300/600], Loss: 0.1363
 Epoch [2/5], Step [400/600], Loss: 0.1438
 Epoch [2/5], Step [500/600], Loss: 0.1520
 Epoch [2/5], Step [600/600], Loss: 0.0439
 Epoch [3/5], Step [100/600], Loss: 0.1012
 Epoch [3/5], Step [200/600], Loss: 0.0243
 Epoch [3/5], Step [300/600], Loss: 0.0556
 Epoch [3/5], Step [400/600], Loss: 0.0636
 Epoch [3/5], Step [500/600], Loss: 0.0264
 Epoch [3/5], Step [600/600], Loss: 0.0689
 Epoch [4/5], Step [100/600], Loss: 0.0561
 Epoch [4/5], Step [200/600], Loss: 0.0216
 Epoch [4/5], Step [300/600], Loss: 0.0193
 Epoch [4/5], Step [400/600], Loss: 0.0708
 Epoch [4/5], Step [500/600], Loss: 0.0672
 Epoch [4/5], Step [600/600], Loss: 0.0781
 Epoch [5/5], Step [100/600], Loss: 0.0491
 Epoch [5/5], Step [200/600], Loss: 0.0177
 Epoch [5/5], Step [300/600], Loss: 0.0081
 Epoch [5/5], Step [400/600], Loss: 0.0087
 Epoch [5/5], Step [500/600], Loss: 0.0087
 Epoch [5/5], Step [600/600], Loss: 0.0145

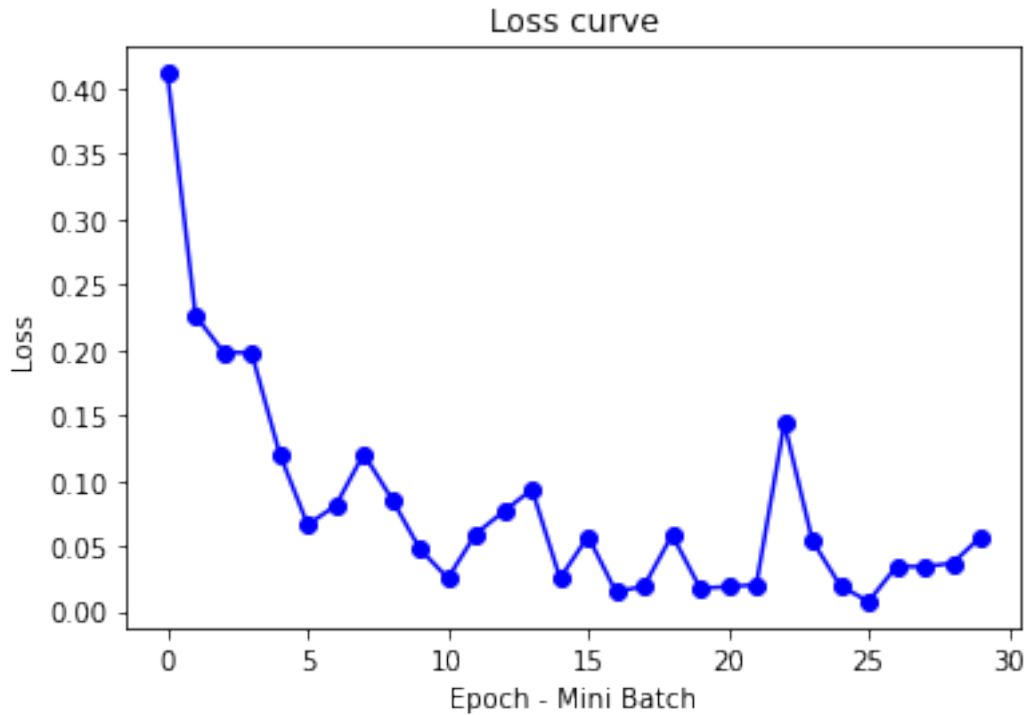


Test Accuracy of the network on 10000 test images: 98.15 %

CNN with Filters = 128 and Hidden layer dimension = 128

#####

Epoch [1/5], Step [100/600], Loss: 0.4118
Epoch [1/5], Step [200/600], Loss: 0.2263
Epoch [1/5], Step [300/600], Loss: 0.1983
Epoch [1/5], Step [400/600], Loss: 0.1977
Epoch [1/5], Step [500/600], Loss: 0.1207
Epoch [1/5], Step [600/600], Loss: 0.0662
Epoch [2/5], Step [100/600], Loss: 0.0809
Epoch [2/5], Step [200/600], Loss: 0.1204
Epoch [2/5], Step [300/600], Loss: 0.0860
Epoch [2/5], Step [400/600], Loss: 0.0477
Epoch [2/5], Step [500/600], Loss: 0.0253
Epoch [2/5], Step [600/600], Loss: 0.0593
Epoch [3/5], Step [100/600], Loss: 0.0765
Epoch [3/5], Step [200/600], Loss: 0.0933
Epoch [3/5], Step [300/600], Loss: 0.0252
Epoch [3/5], Step [400/600], Loss: 0.0573
Epoch [3/5], Step [500/600], Loss: 0.0146
Epoch [3/5], Step [600/600], Loss: 0.0193
Epoch [4/5], Step [100/600], Loss: 0.0584
Epoch [4/5], Step [200/600], Loss: 0.0175
Epoch [4/5], Step [300/600], Loss: 0.0188
Epoch [4/5], Step [400/600], Loss: 0.0202
Epoch [4/5], Step [500/600], Loss: 0.1439
Epoch [4/5], Step [600/600], Loss: 0.0539
Epoch [5/5], Step [100/600], Loss: 0.0193
Epoch [5/5], Step [200/600], Loss: 0.0074
Epoch [5/5], Step [300/600], Loss: 0.0343
Epoch [5/5], Step [400/600], Loss: 0.0341
Epoch [5/5], Step [500/600], Loss: 0.0368
Epoch [5/5], Step [600/600], Loss: 0.0567



Test Accuracy of the network on 10000 test images: 98.08 %

CNN with Filters = 256 and Hidden layer dimension = 128


#####

Epoch [1/5], Step [100/600], Loss: 0.2774
 Epoch [1/5], Step [200/600], Loss: 0.1920
 Epoch [1/5], Step [300/600], Loss: 0.2147
 Epoch [1/5], Step [400/600], Loss: 0.0438
 Epoch [1/5], Step [500/600], Loss: 0.2096
 Epoch [1/5], Step [600/600], Loss: 0.0244
 Epoch [2/5], Step [100/600], Loss: 0.0260
 Epoch [2/5], Step [200/600], Loss: 0.0400
 Epoch [2/5], Step [300/600], Loss: 0.0322
 Epoch [2/5], Step [400/600], Loss: 0.0587
 Epoch [2/5], Step [500/600], Loss: 0.0587
 Epoch [2/5], Step [600/600], Loss: 0.0362
 Epoch [3/5], Step [100/600], Loss: 0.0562
 Epoch [3/5], Step [200/600], Loss: 0.0398
 Epoch [3/5], Step [300/600], Loss: 0.0137
 Epoch [3/5], Step [400/600], Loss: 0.1128
 Epoch [3/5], Step [500/600], Loss: 0.0406
 Epoch [3/5], Step [600/600], Loss: 0.0431
 Epoch [4/5], Step [100/600], Loss: 0.0217

```

Epoch [4/5], Step [200/600], Loss: 0.0188
Epoch [4/5], Step [300/600], Loss: 0.0659
Epoch [4/5], Step [400/600], Loss: 0.0201
Epoch [4/5], Step [500/600], Loss: 0.0041
Epoch [4/5], Step [600/600], Loss: 0.1107
Epoch [5/5], Step [100/600], Loss: 0.0025
Epoch [5/5], Step [200/600], Loss: 0.0193
Epoch [5/5], Step [300/600], Loss: 0.0394
Epoch [5/5], Step [400/600], Loss: 0.0056
Epoch [5/5], Step [500/600], Loss: 0.0753
Epoch [5/5], Step [600/600], Loss: 0.0625

```



```

#### Test Accuracy of the network on 10000 test images: 97.68 % ####

```


```

[22]: for j in filter:
        print("\n\n##### CNN with Filters = ", j, " and Hidden layer_
        ↪dimension = 256 #####")
        model = CNN(j, 3, 256).to(device)
        optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
        train(model, optimizer, "CNN")
        predict(model, type = "CNN")

```

CNN with Filters = 32 and Hidden layer dimension = 256

Epoch [1/5], Step [100/600], Loss: 0.2377
Epoch [1/5], Step [200/600], Loss: 0.3210
Epoch [1/5], Step [300/600], Loss: 0.1523
Epoch [1/5], Step [400/600], Loss: 0.0878
Epoch [1/5], Step [500/600], Loss: 0.1520
Epoch [1/5], Step [600/600], Loss: 0.1249
Epoch [2/5], Step [100/600], Loss: 0.0530
Epoch [2/5], Step [200/600], Loss: 0.0568
Epoch [2/5], Step [300/600], Loss: 0.1605
Epoch [2/5], Step [400/600], Loss: 0.0707
Epoch [2/5], Step [500/600], Loss: 0.0518
Epoch [2/5], Step [600/600], Loss: 0.0603
Epoch [3/5], Step [100/600], Loss: 0.0398
Epoch [3/5], Step [200/600], Loss: 0.0163
Epoch [3/5], Step [300/600], Loss: 0.0161
Epoch [3/5], Step [400/600], Loss: 0.0204
Epoch [3/5], Step [500/600], Loss: 0.0832
Epoch [3/5], Step [600/600], Loss: 0.0342
Epoch [4/5], Step [100/600], Loss: 0.0193
Epoch [4/5], Step [200/600], Loss: 0.0371
Epoch [4/5], Step [300/600], Loss: 0.0122
Epoch [4/5], Step [400/600], Loss: 0.0136
Epoch [4/5], Step [500/600], Loss: 0.0056
Epoch [4/5], Step [600/600], Loss: 0.0138
Epoch [5/5], Step [100/600], Loss: 0.0167
Epoch [5/5], Step [200/600], Loss: 0.0253
Epoch [5/5], Step [300/600], Loss: 0.0074
Epoch [5/5], Step [400/600], Loss: 0.0084
Epoch [5/5], Step [500/600], Loss: 0.0476
Epoch [5/5], Step [600/600], Loss: 0.0221



Test Accuracy of the network on 10000 test images: 98.55 %


CNN with Filters = 64 and Hidden layer dimension = 256

Epoch [1/5], Step [100/600], Loss: 0.3570
 Epoch [1/5], Step [200/600], Loss: 0.0715
 Epoch [1/5], Step [300/600], Loss: 0.1434
 Epoch [1/5], Step [400/600], Loss: 0.0890
 Epoch [1/5], Step [500/600], Loss: 0.1200
 Epoch [1/5], Step [600/600], Loss: 0.1063
 Epoch [2/5], Step [100/600], Loss: 0.0634
 Epoch [2/5], Step [200/600], Loss: 0.0317
 Epoch [2/5], Step [300/600], Loss: 0.0433
 Epoch [2/5], Step [400/600], Loss: 0.0637
 Epoch [2/5], Step [500/600], Loss: 0.0299
 Epoch [2/5], Step [600/600], Loss: 0.0201
 Epoch [3/5], Step [100/600], Loss: 0.0234
 Epoch [3/5], Step [200/600], Loss: 0.0070
 Epoch [3/5], Step [300/600], Loss: 0.0492
 Epoch [3/5], Step [400/600], Loss: 0.0315
 Epoch [3/5], Step [500/600], Loss: 0.0333
 Epoch [3/5], Step [600/600], Loss: 0.0038
 Epoch [4/5], Step [100/600], Loss: 0.0124
 Epoch [4/5], Step [200/600], Loss: 0.0239

```

Epoch [4/5], Step [300/600], Loss: 0.0254
Epoch [4/5], Step [400/600], Loss: 0.0416
Epoch [4/5], Step [500/600], Loss: 0.0023
Epoch [4/5], Step [600/600], Loss: 0.0248
Epoch [5/5], Step [100/600], Loss: 0.0204
Epoch [5/5], Step [200/600], Loss: 0.0039
Epoch [5/5], Step [300/600], Loss: 0.0131
Epoch [5/5], Step [400/600], Loss: 0.0098
Epoch [5/5], Step [500/600], Loss: 0.0308
Epoch [5/5], Step [600/600], Loss: 0.0112

```



Test Accuracy of the network on 10000 test images: 98.36 %


CNN with Filters = 128 and Hidden layer dimension = 256
#####

```

Epoch [1/5], Step [100/600], Loss: 0.2840
Epoch [1/5], Step [200/600], Loss: 0.1665
Epoch [1/5], Step [300/600], Loss: 0.1532
Epoch [1/5], Step [400/600], Loss: 0.0424
Epoch [1/5], Step [500/600], Loss: 0.0532
Epoch [1/5], Step [600/600], Loss: 0.0879
Epoch [2/5], Step [100/600], Loss: 0.1035
Epoch [2/5], Step [200/600], Loss: 0.1803

```

Epoch [2/5], Step [300/600], Loss: 0.0280
 Epoch [2/5], Step [400/600], Loss: 0.1494
 Epoch [2/5], Step [500/600], Loss: 0.1776
 Epoch [2/5], Step [600/600], Loss: 0.0280
 Epoch [3/5], Step [100/600], Loss: 0.0173
 Epoch [3/5], Step [200/600], Loss: 0.0391
 Epoch [3/5], Step [300/600], Loss: 0.0619
 Epoch [3/5], Step [400/600], Loss: 0.0254
 Epoch [3/5], Step [500/600], Loss: 0.0576
 Epoch [3/5], Step [600/600], Loss: 0.0156
 Epoch [4/5], Step [100/600], Loss: 0.0073
 Epoch [4/5], Step [200/600], Loss: 0.0247
 Epoch [4/5], Step [300/600], Loss: 0.0155
 Epoch [4/5], Step [400/600], Loss: 0.0214
 Epoch [4/5], Step [500/600], Loss: 0.0593
 Epoch [4/5], Step [600/600], Loss: 0.0137
 Epoch [5/5], Step [100/600], Loss: 0.0162
 Epoch [5/5], Step [200/600], Loss: 0.0368
 Epoch [5/5], Step [300/600], Loss: 0.0155
 Epoch [5/5], Step [400/600], Loss: 0.0027
 Epoch [5/5], Step [500/600], Loss: 0.0187
 Epoch [5/5], Step [600/600], Loss: 0.0130




Test Accuracy of the network on 10000 test images: 98.53 %

CNN with Filters = 256 and Hidden layer dimension = 256

#####

Epoch [1/5], Step [100/600], Loss: 0.3685
Epoch [1/5], Step [200/600], Loss: 0.2221
Epoch [1/5], Step [300/600], Loss: 0.1338
Epoch [1/5], Step [400/600], Loss: 0.1702
Epoch [1/5], Step [500/600], Loss: 0.1092
Epoch [1/5], Step [600/600], Loss: 0.1561
Epoch [2/5], Step [100/600], Loss: 0.0492
Epoch [2/5], Step [200/600], Loss: 0.0232
Epoch [2/5], Step [300/600], Loss: 0.1327
Epoch [2/5], Step [400/600], Loss: 0.1150
Epoch [2/5], Step [500/600], Loss: 0.0110
Epoch [2/5], Step [600/600], Loss: 0.0422
Epoch [3/5], Step [100/600], Loss: 0.0614
Epoch [3/5], Step [200/600], Loss: 0.0554
Epoch [3/5], Step [300/600], Loss: 0.0554
Epoch [3/5], Step [400/600], Loss: 0.0339
Epoch [3/5], Step [500/600], Loss: 0.0303
Epoch [3/5], Step [600/600], Loss: 0.0219
Epoch [4/5], Step [100/600], Loss: 0.0412
Epoch [4/5], Step [200/600], Loss: 0.0110
Epoch [4/5], Step [300/600], Loss: 0.0158
Epoch [4/5], Step [400/600], Loss: 0.0400
Epoch [4/5], Step [500/600], Loss: 0.0173
Epoch [4/5], Step [600/600], Loss: 0.0098
Epoch [5/5], Step [100/600], Loss: 0.0018
Epoch [5/5], Step [200/600], Loss: 0.0052
Epoch [5/5], Step [300/600], Loss: 0.0447
Epoch [5/5], Step [400/600], Loss: 0.0008
Epoch [5/5], Step [500/600], Loss: 0.0014
Epoch [5/5], Step [600/600], Loss: 0.0156




Test Accuracy of the network on 10000 test images: 98.34 %

```
[23]: for j in filter:
    print("\n\n##### CNN with Filters = ", j, " and Hidden layer_
    ↪dimension = 512 #####")
    model = CNN(j, 3, 512).to(device)
    optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
    train(model, optimizer, "CNN")
    predict(model, type = "CNN")
```

```
##### CNN with Filters = 32 and Hidden layer dimension = 512 #####
Epoch [1/5], Step [100/600], Loss: 0.2476
Epoch [1/5], Step [200/600], Loss: 0.3193
Epoch [1/5], Step [300/600], Loss: 0.0867
Epoch [1/5], Step [400/600], Loss: 0.2249
Epoch [1/5], Step [500/600], Loss: 0.0991
Epoch [1/5], Step [600/600], Loss: 0.0523
Epoch [2/5], Step [100/600], Loss: 0.0659
Epoch [2/5], Step [200/600], Loss: 0.0890
Epoch [2/5], Step [300/600], Loss: 0.1004
Epoch [2/5], Step [400/600], Loss: 0.0325
Epoch [2/5], Step [500/600], Loss: 0.0492
```

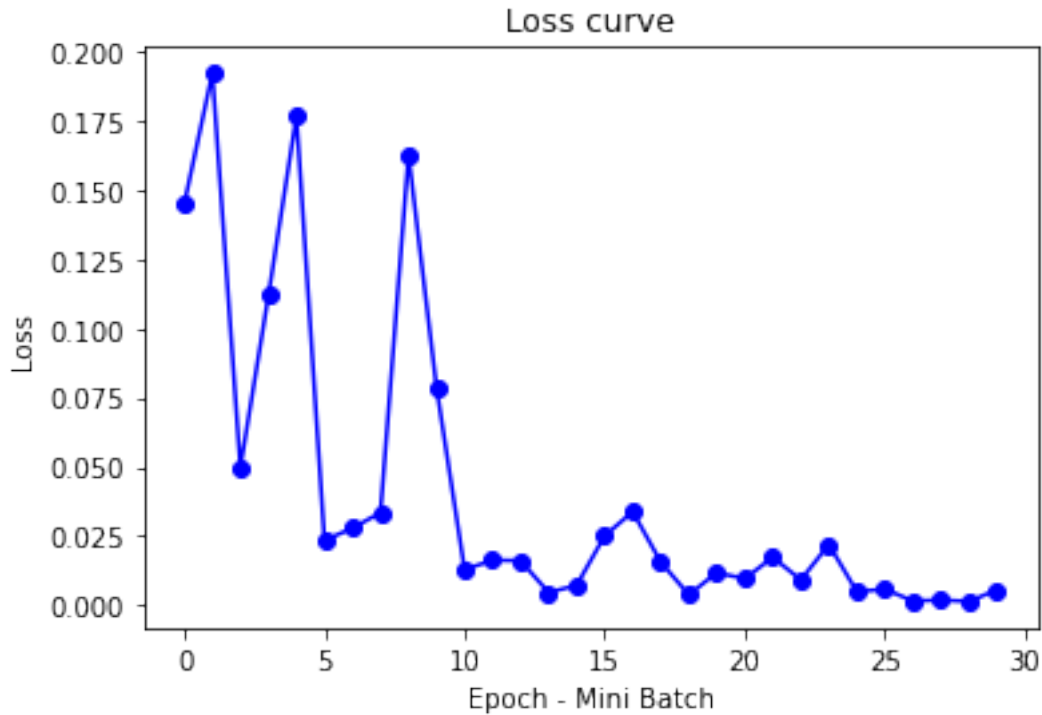
Epoch [2/5], Step [600/600], Loss: 0.0584
 Epoch [3/5], Step [100/600], Loss: 0.0286
 Epoch [3/5], Step [200/600], Loss: 0.0288
 Epoch [3/5], Step [300/600], Loss: 0.0470
 Epoch [3/5], Step [400/600], Loss: 0.0306
 Epoch [3/5], Step [500/600], Loss: 0.0423
 Epoch [3/5], Step [600/600], Loss: 0.0498
 Epoch [4/5], Step [100/600], Loss: 0.0136
 Epoch [4/5], Step [200/600], Loss: 0.0157
 Epoch [4/5], Step [300/600], Loss: 0.0025
 Epoch [4/5], Step [400/600], Loss: 0.0236
 Epoch [4/5], Step [500/600], Loss: 0.0107
 Epoch [4/5], Step [600/600], Loss: 0.0493
 Epoch [5/5], Step [100/600], Loss: 0.0195
 Epoch [5/5], Step [200/600], Loss: 0.0042
 Epoch [5/5], Step [300/600], Loss: 0.0431
 Epoch [5/5], Step [400/600], Loss: 0.0170
 Epoch [5/5], Step [500/600], Loss: 0.0461
 Epoch [5/5], Step [600/600], Loss: 0.0346



Test Accuracy of the network on 10000 test images: 98.55 %

CNN with Filters = 64 and Hidden layer dimension = 512

Epoch [1/5], Step [100/600], Loss: 0.1452
Epoch [1/5], Step [200/600], Loss: 0.1923
Epoch [1/5], Step [300/600], Loss: 0.0498
Epoch [1/5], Step [400/600], Loss: 0.1119
Epoch [1/5], Step [500/600], Loss: 0.1767
Epoch [1/5], Step [600/600], Loss: 0.0234
Epoch [2/5], Step [100/600], Loss: 0.0279
Epoch [2/5], Step [200/600], Loss: 0.0334
Epoch [2/5], Step [300/600], Loss: 0.1628
Epoch [2/5], Step [400/600], Loss: 0.0786
Epoch [2/5], Step [500/600], Loss: 0.0128
Epoch [2/5], Step [600/600], Loss: 0.0163
Epoch [3/5], Step [100/600], Loss: 0.0161
Epoch [3/5], Step [200/600], Loss: 0.0046
Epoch [3/5], Step [300/600], Loss: 0.0072
Epoch [3/5], Step [400/600], Loss: 0.0251
Epoch [3/5], Step [500/600], Loss: 0.0338
Epoch [3/5], Step [600/600], Loss: 0.0156
Epoch [4/5], Step [100/600], Loss: 0.0038
Epoch [4/5], Step [200/600], Loss: 0.0117
Epoch [4/5], Step [300/600], Loss: 0.0096
Epoch [4/5], Step [400/600], Loss: 0.0173
Epoch [4/5], Step [500/600], Loss: 0.0091
Epoch [4/5], Step [600/600], Loss: 0.0220
Epoch [5/5], Step [100/600], Loss: 0.0049
Epoch [5/5], Step [200/600], Loss: 0.0058
Epoch [5/5], Step [300/600], Loss: 0.0015
Epoch [5/5], Step [400/600], Loss: 0.0019
Epoch [5/5], Step [500/600], Loss: 0.0013
Epoch [5/5], Step [600/600], Loss: 0.0056




Test Accuracy of the network on 10000 test images: 98.24 %

CNN with Filters = 128 and Hidden layer dimension = 512

#####

Epoch [1/5], Step [100/600], Loss: 0.3267
 Epoch [1/5], Step [200/600], Loss: 0.1744
 Epoch [1/5], Step [300/600], Loss: 0.1354
 Epoch [1/5], Step [400/600], Loss: 0.0951
 Epoch [1/5], Step [500/600], Loss: 0.0744
 Epoch [1/5], Step [600/600], Loss: 0.1074
 Epoch [2/5], Step [100/600], Loss: 0.0554
 Epoch [2/5], Step [200/600], Loss: 0.0463
 Epoch [2/5], Step [300/600], Loss: 0.0485
 Epoch [2/5], Step [400/600], Loss: 0.0565
 Epoch [2/5], Step [500/600], Loss: 0.0318
 Epoch [2/5], Step [600/600], Loss: 0.0384
 Epoch [3/5], Step [100/600], Loss: 0.0136
 Epoch [3/5], Step [200/600], Loss: 0.0151
 Epoch [3/5], Step [300/600], Loss: 0.0754
 Epoch [3/5], Step [400/600], Loss: 0.0225
 Epoch [3/5], Step [500/600], Loss: 0.0959
 Epoch [3/5], Step [600/600], Loss: 0.0129
 Epoch [4/5], Step [100/600], Loss: 0.0185

Epoch [4/5], Step [200/600], Loss: 0.0147
 Epoch [4/5], Step [300/600], Loss: 0.0053
 Epoch [4/5], Step [400/600], Loss: 0.0662
 Epoch [4/5], Step [500/600], Loss: 0.0077
 Epoch [4/5], Step [600/600], Loss: 0.0140
 Epoch [5/5], Step [100/600], Loss: 0.0502
 Epoch [5/5], Step [200/600], Loss: 0.0036
 Epoch [5/5], Step [300/600], Loss: 0.0074
 Epoch [5/5], Step [400/600], Loss: 0.0044
 Epoch [5/5], Step [500/600], Loss: 0.0205
 Epoch [5/5], Step [600/600], Loss: 0.0248




Test Accuracy of the network on 10000 test images: 98.3 %

CNN with Filters = 256 and Hidden layer dimension = 512

#####

Epoch [1/5], Step [100/600], Loss: 0.3187
 Epoch [1/5], Step [200/600], Loss: 0.1807
 Epoch [1/5], Step [300/600], Loss: 0.1779
 Epoch [1/5], Step [400/600], Loss: 0.1014
 Epoch [1/5], Step [500/600], Loss: 0.0459
 Epoch [1/5], Step [600/600], Loss: 0.1256
 Epoch [2/5], Step [100/600], Loss: 0.0289

Epoch [2/5], Step [200/600], Loss: 0.0342
Epoch [2/5], Step [300/600], Loss: 0.1428
Epoch [2/5], Step [400/600], Loss: 0.0270
Epoch [2/5], Step [500/600], Loss: 0.1256
Epoch [2/5], Step [600/600], Loss: 0.0568
Epoch [3/5], Step [100/600], Loss: 0.0237
Epoch [3/5], Step [200/600], Loss: 0.0252
Epoch [3/5], Step [300/600], Loss: 0.0295
Epoch [3/5], Step [400/600], Loss: 0.0496
Epoch [3/5], Step [500/600], Loss: 0.0573
Epoch [3/5], Step [600/600], Loss: 0.0380
Epoch [4/5], Step [100/600], Loss: 0.0347
Epoch [4/5], Step [200/600], Loss: 0.0072
Epoch [4/5], Step [300/600], Loss: 0.0597
Epoch [4/5], Step [400/600], Loss: 0.0122
Epoch [4/5], Step [500/600], Loss: 0.0417
Epoch [4/5], Step [600/600], Loss: 0.0542
Epoch [5/5], Step [100/600], Loss: 0.0043
Epoch [5/5], Step [200/600], Loss: 0.0093
Epoch [5/5], Step [300/600], Loss: 0.0206
Epoch [5/5], Step [400/600], Loss: 0.0022
Epoch [5/5], Step [500/600], Loss: 0.0433
Epoch [5/5], Step [600/600], Loss: 0.0179



Test Accuracy of the network on 10000 test images: 98.16 %

- 1.1.2 In general, from the plots above, it is noted that, the Test accuracy decreases (in most cases) with increase in filter sizes (from 32 to 256)
- 1.1.3 In this case, the hidden layer dimensions = 256 performed better in all choices of filter sizes than the models with hidden layer dimension = 128 and 512. And the model with 512 hidden layer dimension performed better than the model with 128 hidden layer dimension.